

# The Spark Foundation

**Predict the percentage of an student based on the no. of study hours.**

**ML Algorithm = Simple Linear Regression**

**Problem To Be Solved = What will be predicted score if a student studies for 9.25 hrs/ day? ¶**

```
In [1]: #importing all the libraries.  
import numpy as np  
import matplotlib.pyplot as plt  
import pandas as pd  
import seaborn as sns
```

```
In [2]: #importing and reaadng the dataset  
dataset = pd.read_csv("student_scores.csv")
```

```
In [3]: #the first five values in the dataset  
dataset.head()
```

Out[3]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [4]: #number of rows and columns  
dataset.shape
```

Out[4]: (25, 2)

```
In [5]: dataset.describe()
```

Out[5]:


	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [6]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    Hours    25 non-null    float64
1    Scores   25 non-null    int64   
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```


## Visualization Of Data

```
In [7]: #Hours Vs Percentage of Scores
plt.figure(figsize=(10,6))
plt.scatter(dataset['Hours'], dataset['Scores'], c='black')
plt.title('Hours vs Percentage')
plt.xlabel('Hours Studied')
plt.ylabel('Scores')
plt.show()
```




```
In [8]: plt.figure(figsize=(10,6))  
plt.bar(dataset['Hours'], dataset['Scores'], color='blue')  
plt.xlabel('Hours Studied')  
plt.ylabel('Scores')
```

Out[8]: Text(0, 0.5, 'Scores')




```
In [9]: plt.figure(figsize=(15,6))  
sns.barplot(x='Hours', y='Scores', data=dataset)
```

Out[9]: <AxesSubplot:xlabel='Hours', ylabel='Scores'>



```
In [10]: plt.figure(figsize=(10,6))  
sns.lineplot(data=dataset)
```

Out[10]: <AxesSubplot:>



## Train-Test separation of data

```
In [11]: #X will take all the values except for the last column which is our dependent variable  
X = dataset.iloc[:, :-1].values  
y = dataset.iloc[:, -1].values
```

```
In [12]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```


## Training the Simple Linear Regression model on the Training set

```
In [13]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

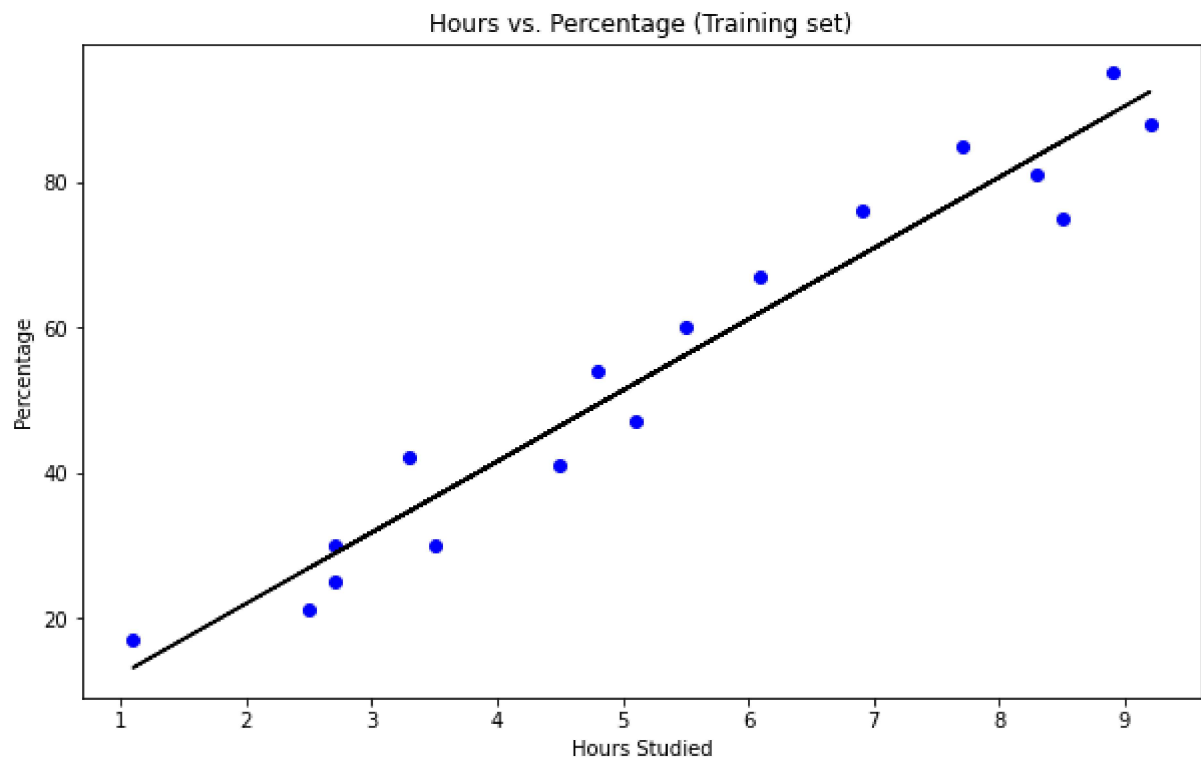
Out[13]: LinearRegression()

```
In [14]: # Plotting the regression line
line = regressor.coef_*X+regressor.intercept_

# Plotting for the test data
plt.figure(figsize=(10,6))
plt.scatter(X, y, c='coral')
plt.plot(X, line,c = 'black');
plt.show()
```



```
In [15]: #Visualising the Training set results
plt.figure(figsize=(10,6))
plt.scatter(X_train, y_train, c='blue')
plt.plot(X_train, regressor.predict(X_train), c='black')
plt.title('Hours vs. Percentage (Training set)')
plt.xlabel('Hours Studied')
plt.ylabel('Percentage')
plt.show()
```




In [16]: *#Predicting the Test set results*

```
y_pred = regressor.predict(X_test)
print(y_pred)
```

```
[17.05366541 33.69422878 74.80620886 26.8422321  60.12335883 39.56736879
 20.96909209 78.72163554]
```

In [17]: *#Visualising the Test set results*

```
plt.figure(figsize=(10,6))
plt.scatter(X_test, y_test, color = 'green')
plt.plot(X_train, regressor.predict(X_train), color = 'black')
plt.title('Hours vs. Percentage (Test set)')
plt.xlabel('Hours Studied')
plt.ylabel('Percentage')
plt.show()
```






```
In [18]: #Comparing the actual values with the predicted ones.
ds = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
ds
```

Out[18]:

	Actual	Predicted
0	20	17.053665
1	27	33.694229
2	69	74.806209
3	30	26.842232
4	62	60.123359
5	35	39.567369
6	24	20.969092
7	86	78.721636

```
In [19]: plt.figure(figsize=(8,5))
sns.heatmap(ds.corr())
```

Out[19]: <AxesSubplot:>



## Predicting the score if the student studies for 9.25 hours/day

```
In [20]: dataset = np.array(9.25)
dataset = dataset.reshape(-1, 1)
pred = regressor.predict(dataset)
print("If the student studies for 9.25 hours/day, the score is {}".format(pred))
```

If the student studies for 9.25 hours/day, the score is [92.91505723].

## Check the predicting score if the student studies for 6 hours/day

```
In [21]: dataset = np.array(6)
dataset = dataset.reshape(-1, 1)
pred = regressor.predict(dataset)
print("If the student studies for 6 hours/day, the score is {}".format(pred))
```

If the student studies for 6 hours/day, the score is [61.1022155].

## Conclusion:

**We used a Linear Regression Model to predict the score of a student if he/she studies for 9.25 hours/day and the Predicted Score came out to be 92.91.**